Clarifying Questions

Data & Scope

- What time range do you want to include initially, and should we backfill or only move forward?
- Which source systems feed this (ED tracker, EHR, patient experience tool, staffing/rostering, bed management)?
- What is the canonical timestamp for "arrival," "triage," and "first clinical contact"?
- Do we track acuity (e.g., triage level) and chief complaint? If yes, which fields?
- Which encounters should be excluded (test patients, incomplete registrations, inter-facility transfers, LWBS, etc.)?

Metrics & Definitions

- How do you define "Seen within 30 minutes" (door→triage, door→provider, registration→provider)?
- Is "Wait Time" mean, median, or both? Do we cap outliers?
- What counts as "Admission" (obs included or separate)? What counts as a "Referral"?
- How is Patient Satisfaction captured (scale, timing, response window)? Any minimum N to display?

Refresh & Performance

- How often should the dashboard refresh (near-real-time, hourly, daily)?
- Do we need row-level security (e.g., by unit or role) and any data masking?
- Any performance or capacity constraints for the Power BI workspace?

Audience & UX

- Who are the primary users (charge nurse, ED director, bed manager, execs)? What are their top 3 questions?
- Preferred layout (single-page overview vs. multi-page journey)? Any must-have visuals?
- Accessibility needs (high contrast, font sizes, color-blind palette, tooltips in plain language)?

Governance & Rollout

- Who owns metric definitions and change control?
- Do we need an audit trail / data dictionary embedded in the report?
- What's the approval path for publishing and sharing?

Assumptions

- Data is de-identified (or appropriately authorized) and safe for aggregate reporting.
- Timestamps are recorded consistently across systems and in a single time zone.
- The intake process (registration → triage → provider) is broadly unchanged over the period analyzed.
- "Admission" and "Observation" are coded consistently; referral destinations use a stable naming convention.
- Satisfaction scores are comparable across time and collection methods.
- Missing or obviously erroneous values will be excluded via simple rules (e.g., negative durations, extreme outliers).
- We will begin with descriptive analytics (no forecasting or acuity risk modeling in v1).
- Power BI Pro (or Premium) capacity is available, and users have access to the workspace/app.
- The design will use a calm blue theme, clear legends, and consistent number formatting.
- Stakeholders will provide feedback in short, iterative cycles (e.g., biweekly).

Caveats

- Descriptive, not causal: the dashboard shows patterns, not causes; operational trials are needed to prove impact.
- Data completeness & bias: survey response bias and missing acuity can skew experience and flow insights.

- Definition sensitivity: small changes to metric definitions (e.g., what “first contact” means) can shift results materially.
- Seasonality & anomalies: unusual events (holidays, outbreaks, system downtimes) can distort trends.
- Downstream dependencies: improvements may be limited by inpatient bed availability and consult response times.
- Latency: if refresh is not real-time, the view may lag current conditions.
- Privacy & governance: any drill-through to row-level data must follow privacy rules and role-based access.